

Assembly Code Analysis: GD vs Predicted

1. C Code to Assembly Mapping (GD Version)

Function Signature and Register Mapping

```
c
int *func0(const int numbers[], int size, int delimiter, int *out_size)
```

Register Assignments:

- (x0) (x20) = (numbers[]) (input array pointer)
- (x1) (x21) = (size) (array size)
- (x2) (x19) = (delimiter) (delimiter value)
- (x3) = (out_size) (pointer to output size)

Line-by-Line C to Assembly Mapping

C Line: `*out_size = size > 0 ? (size * 2) - 1 : 0;`

```
assembly

lsl w8, w1, #1      ; w8 = size * 2 (left shift by 1)
sub w8, w8, #1      ; w8 = (size * 2) - 1
cmp w1, #0          ; compare size with 0
csel w8, w8, wzr, gt ; if size > 0, use w8, else use 0
str w8, [x3]        ; store result in *out_size
```

C Line: `int *out = (int *)malloc(*out_size * sizeof(int));`

```
assembly

sbfiz x0, x8, #2, #32 ; x0 = w8 * 4 (sizeof(int))
bl _malloc            ; call malloc, result in x0
```

C Line: `if (size > 0) out[0] = numbers[0];`

```
assembly

cmp w21, #1          ; compare size with 1
b.lt LBB0_9          ; if size < 1, jump to end
ldr w8, [x20]         ; load numbers[0]
str w8, [x0]          ; store in out[0]
b.eq LBB0_9          ; if size == 1, jump to end
```

C Line: Loop `for (int i = 1, j = 1; i < size; ++i)`

The vectorized loop (LBB0_4-LBB0_6) processes multiple elements at once:

```
assembly
; Vectorized processing of 16 elements at a time
dup.4s v0, w19          ; duplicate delimiter into vector
st2.4s { v0, v1 }, [x16] ; store interleaved delimiter and data
```

2. Vertical Comparison: GD vs Predicted

Key Differences Identified

Line	GD Version	Predicted Version	Difference Type
27	<code>b.eq LBB0_9</code>	<code>cmp w21, #1</code> + <code>b.eq LBB0_9</code>	CRITICAL: Redundant comparison
30	<code>sub x11, x8, #1</code>	<code>sub x10, x8, #1</code>	Register name
31	<code>cmp x11, #16</code>	<code>cmp x10, #8</code>	CRITICAL: Different vectorization factor
33	<code>mov w10, #1</code>	<code>mov w11, #1</code>	Register name
36	<code>bfi x10, x13, #5, #59</code>	<code>bfi x11, x12, #1, #63</code>	CRITICAL: Different bit field insertion
38-56	16-element vectorization	8-element vectorization	CRITICAL: Different SIMD strategy
63	<code>stp w19, w10, [x9, #-4]</code>	<code>stur w19, [x9, #-4]</code> + <code>str w10, [x9], #8</code>	CRITICAL: Different store pattern

3. Critical Error Analysis

Error 1: Redundant Comparison (Line 27)

GD (Correct):

```
assembly
str w8, [x0]          ; store numbers[0] in out[0]
b.eq LBB0_9           ; if size == 1, exit (condition already set)
```

Predicted (Incorrect):

assembly

```
str w8, [x0]           ; store numbers[0] in out[0]
cmp w21, #1            ; REDUNDANT: compare size with 1 again
b.eq LBB0_9            ; if size == 1, exit
```

Problem: The predicted version performs a redundant comparison. The condition was already set by the previous `cmp w21, #1` instruction.

Error 2: Incorrect Vectorization Factor

GD (Correct): Processes 16 elements per iteration

assembly

```
cmp x11, #16           ; check if remaining >= 16
and x12, x11, #0xfffffffffff0 ; align to 16-element boundary
```

Predicted (Incorrect): Processes 8 elements per iteration

assembly

```
cmp x10, #8            ; check if remaining >= 8
and x12, x10, #0xfffffffffff8 ; align to 8-element boundary
```

Error 3: Incorrect Bit Field Insertion

GD (Correct):

assembly

```
bfi x10, x13, #5, #59   ; insert bits for 16-element indexing
```

Predicted (Incorrect):

assembly

```
bfi x11, x12, #1, #63   ; insert bits for 8-element indexing
```

Problem: The bit field parameters don't match the vectorization strategy.

Error 4: Incorrect Store Pattern in Scalar Loop

GD (Correct):

assembly

```
stp w19, w10, [x9, #-4] ; store delimiter and number as pair
```

Predicted (Incorrect):

assembly

```
stur w19, [x9, #-4]      ; store delimiter  
str w10, [x9], #8        ; store number with post-increment
```

Problem: The predicted version uses different addressing modes and increment patterns.

4. Connection to O2 Optimization

The errors in the predicted version stem from **incorrect assumptions about compiler optimization strategies**:

Vectorization Strategy Mismatch

- **O2 optimization** typically chooses vectorization factors based on:
 - Target architecture capabilities
 - Data alignment requirements
 - Loop trip count analysis
- The prediction incorrectly assumed 8-element SIMD instead of 16-element

Instruction Scheduling Errors

- **O2 optimization** eliminates redundant comparisons through sophisticated control flow analysis
- The prediction failed to recognize that condition codes were already set

Memory Access Pattern Optimization

- **O2 optimization** uses paired stores (`((stp))`) for better memory bandwidth utilization
- The prediction used separate stores, which is less efficient

5. Summary

The predicted assembly contains **4 critical logical errors** that would cause incorrect program behavior:

1. **Redundant comparison** - wastes cycles
2. **Wrong vectorization factor** - processes fewer elements efficiently
3. **Incorrect bit manipulation** - misaligned memory access patterns
4. **Suboptimal store pattern** - reduced memory throughput

These errors demonstrate the complexity of modern compiler optimizations, where small changes in instruction selection or scheduling can significantly impact both correctness and performance. The O2 optimization level performs sophisticated analysis that's difficult to predict without deep understanding of the compiler's internal algorithms.